

# ZURAIZ ANJUM

AI Engineer | Data Scientist

zuraizwork@gmail.com | +92 306-6354444 | Lahore, Pakistan

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

## PROFILE

AI Engineer with production experience across LLM systems, backend engineering, and applied ML — from fine-tuning and agentic pipelines to full-stack SaaS platforms. Comfortable across the full applied-AI stack: prompt engineering, RAG, model evaluation, API/database integration, and deployment. Fast at picking up new tools and specs, with a track record of shipping tested, production-grade systems.

## TECHNICAL SKILLS

**Languages:** Python, C++, C#, TypeScript, SQL

**AI / ML:** NLP, PyTorch, Scikit-learn, TensorFlow, Keras, NumPy, Pandas

**GenAI & LLMs:** LangChain, LangGraph, Multi-Agent Orchestration, RAG, Prompt Engineering & Evaluation, Fine-tuning, LLaMA, OpenAI & Anthropic SDKs

**Backend & APIs:** FastAPI, Node.js, REST APIs, SSE Streaming, MySQL, PostgreSQL, SQLite

**Tools & Deployment:** Git/GitHub, Docker, Railway, Automated Testing (unit/integration/eval)

**Frontend:** React.js, Next.js, React Native, TypeScript, Tailwind CSS

## PROFESSIONAL EXPERIENCE

**RL Engineer** | *Tensium · UK · Remote* | [tensium.co.uk](https://tensium.co.uk)

**Jul 2026 – Present**

- Building and evaluating reinforcement learning environments and reward pipelines used to train and benchmark agentic model behavior.
- Working directly with the core engineering team on environment design, task specification, and trajectory evaluation for RL training runs.

**Software Engineer** | *Mercor · Contract · Remote* | [mercor.com](https://mercor.com)

**Nov 2025 – Jun 2026**

- Collaborated with a global AI engineering team, designing and optimizing database schemas, APIs, and backend services for AI model integration into ML applications.
- Ran model evaluation to assess accuracy, performance, and reliability of AI outputs; used coding agents and automation to speed up code generation and testing.
- Built production-grade AI infrastructure using Python, SQL, and FastAPI.

**AI/ML Intern** | *Project19 · Remote (US-based)* | [theproject19.com](https://theproject19.com)

**Jun 5 – Aug 5, 2025**

- Fine-tuned and evaluated LLMs for applied use cases; built and tested agentic workflows automating multi-step tasks alongside a remote engineering team on production-facing AI features.

## KEY PROJECTS

**ARES — Autonomous Research & Evidence System** | *Python · LangGraph · FastAPI · RAG · Docker · Railway*

- Designed and deployed a production-grade 20-agent LangGraph research assistant turning natural-language queries into structured research reports, academic papers, visualizations, and comparison matrices.
- Built multi-file RAG (PDF/DOCX/TXT) with ChromaDB + FlashRank reranking, persistent entity memory, and real-time web research via the Tavily API.
- Implemented multi-provider LLM fallback (Groq, Cerebras, Gemini, OpenRouter) with automatic failover, cost/budget tracking, and self-recalibrating confidence scoring.
- Wrote 122 automated tests (unit, integration, groundedness/citation checks) with a critic-agent revision loop enforcing quality thresholds; shipped as a Dockerized REST + SSE streaming API on Railway with Slack/email integrations and LangSmith tracing.

Live demo: [web-production-41acc.up.railway.app](https://web-production-41acc.up.railway.app)

**AI Agent SaaS Platform** | *Next.js · TypeScript · Groq · Stream SDK · Drizzle ORM · Neon DB · Stripe*

- Built a platform where users create AI agents that join real-time video meetings with live chat, LLM-powered responses, and automated post-meeting summaries/transcripts.
- Added Stripe-based subscription billing, auth (email + OAuth), protected routes, and agent/meeting dashboards with history and video playback.

**Crisis Detection Benchmarking & ML Fitness Tracker** | *NLP · DistilBERT/MentalBERT/Flan-T5 · TF-IDF · Scikit-learn*

- Benchmarked DistilBERT, MentalBERT, and Flan-T5 against a TF-IDF + Logistic Regression baseline for crisis detection on social media data, analyzing how closely small models approximate large-scale performance.
- Engineered features from accelerometer/gyroscope data and trained SVM/Random Forest/Neural Network classifiers for exercise recognition, plus a custom rep-counting algorithm.

## LEADERSHIP & ACTIVITIES

Event Coordinator, [SOFTEC](#) — large-scale student-run tech festival, FAST-NUCES.

**Jul 2026 – Jun 2027**

SVP (Senior Vice President), [NUCES Media Group](#), FAST-NUCES.

**Jul 2026 – Jun 2027**

## EDUCATION

BS Data Science — FAST – National University of Computer and Emerging Sciences, Lahore

**2023 – 2027**